

# Rebooting Supercompilation for Haskell

Ömer S. Ağacan  
oagacan@indiana.edu

Ryan R. Newton  
rrnewton@indiana.edu

August 23, 2015

# Rebooting Supercompilation for Haskell - Talk outline

- An overview of supercompilation.

# Rebooting Supercompilation for Haskell - Talk outline

- An overview of supercompilation.
- What's interesting about it in the context of Haskell? Current state-of-the-art.

# Rebooting Supercompilation for Haskell - Talk outline

- An overview of supercompilation.
- What's interesting about it in the context of Haskell? Current state-of-the-art.
- Overview of how it works.

# Rebooting Supercompilation for Haskell - Talk outline

- An overview of supercompilation.
- What's interesting about it in the context of Haskell? Current state-of-the-art.
- Overview of how it works.
- "But where's my supercompiler for Haskell?" My preliminary work and research goals.

# Supercompilation: An overview

- The paper that describes the idea in English: "The Concept of a Supercompiler" Turchin [1986].

# Supercompilation: An overview

- The paper that describes the idea in English: "The Concept of a Supercompiler" Turchin [1986].
- High-level idea:

# Supercompilation: An overview

- The paper that describes the idea in English: "The Concept of a Supercompiler" Turchin [1986].
- High-level idea:
  - Evaluate programs in compile-time, while making the most out of known inputs and definitions.



# Supercompilation: An overview

- The paper that describes the idea in English: "The Concept of a Supercompiler" Turchin [1986].
- High-level idea:
  - Evaluate programs in compile-time, while making the most out of known inputs and definitions.
    - Definitions of used functions.

# Supercompilation: An overview

- The paper that describes the idea in English: "The Concept of a Supercompiler" Turchin [1986].
- High-level idea:
  - Evaluate programs in compile-time, while making the most out of known inputs and definitions.
    - Definitions of used functions.
    - Statically known arguments of functions.

# Supercompilation: An overview

- The paper that describes the idea in English: "The Concept of a Supercompiler" Turchin [1986].
- High-level idea:
  - Evaluate programs in compile-time, while making the most out of known inputs and definitions.
    - Definitions of used functions.
    - Statically known arguments of functions.
    - When branching, propagate learned information through branches and make use of that information while compiling branches. (case expressions)

# Supercompilation: An overview

- The paper that describes the idea in English: "The Concept of a Supercompiler" Turchin [1986].
- High-level idea: (contd)
  - Evaluate programs in compile-time, while making the most out of known inputs and definitions.
    - Most of the time the goal is to generate more efficient programs.  
(but see Klyuchnikov and Romanenko [2010] for a different use of supercompilation)

# Supercompilation in the context of Haskell

- Why is it interesting?
- In a sense, it's the "ultimate" optimization. ("-O99")
- This optimizes in the sense that:  
If we have a programs  $\mathcal{P}_1$  and  $\mathcal{P}_2$ , and  
 $\mathcal{P}_1 \Downarrow v$  in  $N$  steps and  
 $\mathcal{P}_2 \Downarrow v$  in  $M$  steps,  
we consider  $\mathcal{P}_2$  optimized if  $M < N$ .
- An approximation, but works well in practice.

# Supercompilation in the context of Haskell

- It generalizes:
  - Deforestation(Wadler [1988])
  - Partial evaluation
  - Call-pattern specialization(Peyton Jones [2007])
  - Ad-hoc optimizations via rewrite rules, e.g. shortcut fusion (Gill et al. [1993]) or library-specific rewrite rules
  - "Optimizing SYB is Easy!" (Adams et al. [2014]) and "Optimizing Generics is Easy!" (Magalhães et al. [2010]) style "domain-specific" partial evaluators
  - Function specialization(SPECIALIZE pragmas)
  - ... and many more

# Current state-of-the-art

- Bolingbroke [2013] shows some great potential:
  - Up to  $-95.1\%$  runtime improvement.
  - Up to  $-100.0\%$  allocation improvement.
- But it also suffers from problems that are inherent to supercompilation:
  - "We do not attempt to supercompile the full Nofib suite because the other Nofib benchmarks are considerably more complicated and generally suffer from extremely long supercompilation times."  
(Jonsson [201?] focuses on compilation performance, and reports  $<3$  seconds for all the small programs from Nofib)
  - Up to  $+132002.0\%$  in compile time.
  - Up to  $+188.9\%$  in generated code size.

# How it works? An overview

Bolingbroke [2013] laid out a great framework for supercompiling Haskell:

- **Driving:** Take steps according to operational semantics. Some additional steps like case-of-case transformation (Jones and Santos [1998]).



# How it works? An overview

Bolingbroke [2013] laid out a great framework for supercompiling Haskell:

- **Driving:** Take steps according to operational semantics. Some additional steps like case-of-case transformation (Jones and Santos [1998]).
- **Splitting:** When stuck, keep evaluating sub-expressions. Propagate information. After evaluating sub-expressions combine results.

# How it works? An overview

Bolingbroke [2013] laid out a great framework for supercompiling Haskell:

- **Driving:** Take steps according to operational semantics. Some additional steps like case-of-case transformation (Jones and Santos [1998]).
- **Splitting:** When stuck, keep evaluating sub-expressions. Propagate information. After evaluating sub-expressions combine results.
- **Matching:** Evaluating open terms lead to loops. Matcher tried to detect loops, returns information about how to refer to this new loop.

# How it works? An overview

Bolingbroke [2013] laid out a great framework for supercompiling Haskell:

- **Driving:** Take steps according to operational semantics. Some additional steps like case-of-case transformation (Jones and Santos [1998]).
- **Splitting:** When stuck, keep evaluating sub-expressions. Propagate information. After evaluating sub-expressions combine results.
- **Matching:** Evaluating open terms lead to loops. Matcher tried to detect loops, returns information about how to refer to this new loop.
- **Termination checking:** Because perfect matcher is not possible, and some programs just loop.

## How it works? An overview

```
mapOfMap f g = (.) (map f) (map g)
```

```
h1 f g a = map f (map g a)
```

```
h2 f g a =  
  (\f lst -> case lst of  
    []      -> []  
    h : t   -> f h : map f t) f (map g a)
```

```
h3 f g a =  
  case (map g a) of  
    []      -> []  
    h : t   -> f h : map f t
```

## How it works? An overview

```
h3 f g a =  
  case (map g a) of  
    []      -> []  
    h : t -> f h : map f t
```

```
h4 f g a =  
  case (case a of  
    []      -> []  
    h1 : t1 -> g h1 : map g t1) of  
    []      -> []  
    h0 : t0 -> f h0 : map f t0
```

## How it works? An overview

```
h4 f g a =  
  case (case a of  
        []      -> []  
        h1 : t1 -> g h1 : map g t1) of  
    []      -> []  
    h0 : t0 -> f h0 : map f t0
```

Case-of-case transformation: (Jones and Santos [1998])

```
h5 f g a =  
  case a of  
    [] -> case [] of  
        []      -> []  
        h0 : t0 -> f h0 : map f t0  
    h1 : t1 ->  
      case (g h1 : map g t1) of  
        []      -> []  
        h0 : t0 -> f h0 : map f t0
```

## How it works? An overview

```
h5 f g a =  
  case a of  
    [] -> case [] of  
      [] -> []  
      h0 : t0 -> f h0 : map f t0  
    h1 : t1 ->  
      case (g h1 : map g t1) of  
        [] -> []  
        h0 : t0 -> f h0 : map f t0
```

```
h6 f g a =  
  case a of  
    [] -> []  
    h : t -> f (g h) : map f (map g t)
```

## How it works? An overview

```
h6 f g a =  
  case a of  
    []      -> []  
    h : t  -> f (g h) : map f (map g t)
```



## How it works? An overview

```
h6 f g a =  
  case a of  
    []      -> []  
    h : t   -> f (g h) : map f (map g t)
```

---

map f (map g t)

Looks similar to:

```
h1 f g a = map f (map g a)
```

## How it works? An overview

At this point splitter tell us there's a loop.

```
h7 f g a =  
  case a of  
    []      -> []  
    h : t -> f (g h) : h7 f g t
```

Supercompiled version doesn't generate intermediate list.

## Another example, growing arguments

```
reverse_acc []      acc = acc
reverse_acc (h : t) acc = reverse_acc t (h : acc)
goal lst = reverse_acc (reverse_acc lst []) []

h0 lst = reverse_acc (reverse_acc lst []) []
...
h5 lst = case lst of
  []      -> []
  h1 : t1 -> case (reverse_acc t1 (h1 : [])) of
    []      -> []
    h0 : t0 -> reverse_acc t0 (h0 : [])
...
h_ lst = ... reverse_acc t1 (h1 : []) ...
...
h_ lst = ... reverse_acc t2 (h2 : h1 : []) ...
...
```

## Another example, growing arguments

What to do after stopping is completely different story.

Generalization

Rollback

Other/new ideas?

# Problems with supercompilation operations

Each operation has hard problems to solve.

# Problems with supercompilation operations

Each operation has hard problems to solve.

# Problems with supercompilation operations

Each operation has hard problems to solve.

**Splitter:** Propagating too much information may lead to work duplication. Propagating too little information may lead to missing optimization opportunities.

# Problems with supercompilation operations

Each operation has hard problems to solve.

**Splitter:** Propagating too much information may lead to work duplication. Propagating too little information may lead to missing optimization opportunities.

**Matcher:** Matching may lead to work sharing, which may increase memory residency. (Chitil [1997])



# Problems with supercompilation operations

Each operation has hard problems to solve.

**Splitter:** Propagating too much information may lead to work duplication. Propagating too little information may lead to missing optimization opportunities.

**Matcher:** Matching may lead to work sharing, which may increase memory residency. (Chitil [1997])

**Termination checker:** Need to be careful with growing arguments.

# Problems with supercompilation operations

Each operation has hard problems to solve.

**Splitter:** Propagating too much information may lead to work duplication. Propagating too little information may lead to missing optimization opportunities.

**Matcher:** Matching may lead to work sharing, which may increase memory residency. (Chitil [1997])

**Termination checker:** Need to be careful with growing arguments.

Bolingbroke [2013] has some solutions, and documents it nicely.

# Problems with supercompilation operations

Each operation has hard problems to solve.

**Splitter:** Propagating too much information may lead to work duplication. Propagating too little information may lead to missing optimization opportunities.

**Matcher:** Matching may lead to work sharing, which may increase memory residency. (Chitil [1997])

**Termination checker:** Need to be careful with growing arguments.

Bolingbroke [2013] has some solutions, and documents it nicely.

But we still don't have a working implementation.

# "Where's my supercompiler for Haskell?"

- I'm rebooting the supercompiler!

# "Where's my supercompiler for Haskell?"

- I'm rebooting the supercompiler!
- The goal here is to distribute it as a package, downloadable from Hackage.

# "Where's my supercompiler for Haskell?"

- I'm rebooting the supercompiler!
- The goal here is to distribute it as a package, downloadable from Hackage.
- Then the research will follow.

# Current status and problems

There has been some changes in GHC:

- Some changes in the Core theory: Roles.
- Lots of refactoring.

GHC API related problems:

- Some needed internals are not exposed by GHC – requires some modifications in GHC.
- No easy ways to do most basic stuff: Moving terms around(substitutions), known-case reduction, case-of-case, etc. (all done in some parts of Core-to-Core passes, need to reverse engineer)
- No easy way to annotate Core syntax. Duplicating the syntax means duplicating huge amounts of code.
- Working on Core is hard: Invariants are encoded as partial functions without any helpful error messages – if we're lucky, there's a NOTE.
- Some things are not clear. (Types are first-class, but can I use them wherever I want? The Core definition allows this)

# Roadmap

Have a working implementation of supercompiler described in Bolingbroke [2013].



# Roadmap

Have a working implementation of supercompiler described in Bolingbroke [2013].

(Maybe) Improve GHC plugin API on the way.

# Roadmap

Have a working implementation of supercompiler described in Bolingbroke [2013].

(Maybe) Improve GHC plugin API on the way.

Collecting benchmark programs - send yours! (with expected optimizations)

We can collect something like Nofib, but for supercompilation related problems.

# Roadmap

Have a working implementation of supercompiler described in Bolingbroke [2013].

(Maybe) Improve GHC plugin API on the way.

Collecting benchmark programs - send yours! (with expected optimizations)

We can collect something like Nofib, but for supercompilation related problems.

Once we have a working implementation:

# Roadmap

Have a working implementation of supercompiler described in Bolingbroke [2013].

(Maybe) Improve GHC plugin API on the way.

Collecting benchmark programs - send yours! (with expected optimizations)

We can collect something like Nofib, but for supercompilation related problems.

Once we have a working implementation:

- Focus on specific parts(matcher, splitter etc.). Try other ideas from the literature(e.g. homeomorphic embedding).

# Roadmap

Have a working implementation of supercompiler described in Bolingbroke [2013].

(Maybe) Improve GHC plugin API on the way.

Collecting benchmark programs - send yours! (with expected optimizations)

We can collect something like Nofib, but for supercompilation related problems.

Once we have a working implementation:

- Focus on specific parts(matcher, splitter etc.). Try other ideas from the literature(e.g. homeomorphic embedding).
- Work on some of the obvious improvements, like parallelizing the matcher.

# Roadmap

Have a working implementation of supercompiler described in Bolingbroke [2013].

(Maybe) Improve GHC plugin API on the way.

Collecting benchmark programs - send yours! (with expected optimizations)

We can collect something like Nofib, but for supercompilation related problems.

Once we have a working implementation:

- Focus on specific parts(matcher, splitter etc.). Try other ideas from the literature(e.g. homeomorphic embedding).
- Work on some of the obvious improvements, like parallelizing the matcher.
- I'm open for more ideas!

# Thanks!

Github: [osa1/sc-plugin](#)

IRC: [osa1](#)

Email me your slow programs: [oagacan@indiana.edu](mailto:oagacan@indiana.edu)

# References I

- M. D. Adams, A. Farmer, and J. P. Magalhães. Optimizing SYB is Easy! In *Proceedings of the ACM SIGPLAN 2014 Workshop on Partial Evaluation and Program Manipulation*, PEPM '14, New York, NY, USA, 2014. ACM. ISBN 978-1-4503-2619-3. doi: 10.1145/2543728.2543730.
- M. C. Bolingbroke. Call-by-need supercompilation. Technical Report UCAM-CL-TR-835, University of Cambridge, Computer Laboratory, May 2013. URL <http://www.cl.cam.ac.uk/techreports/UCAM-CL-TR-835.pdf>.
- O. Chitil. Common Subexpression Elimination in a Lazy Functional Language, 1997.



## References II

- A. Gill, J. Launchbury, and S. L. Peyton Jones. A Short Cut to Deforestation. In *Proceedings of the Conference on Functional Programming Languages and Computer Architecture*, FPCA '93, pages 223–232, New York, NY, USA, 1993. ACM. ISBN 0-89791-595-X. doi: 10.1145/165180.165214. URL <http://doi.acm.org/10.1145/165180.165214>.
- S. L. P. Jones and A. L. Santos. A transformation-based optimiser for Haskell. In *SCIENCE OF COMPUTER PROGRAMMING*, pages 3–47. Elsevier North-Holland, Inc., 1998.
- P. A. Jonsson. Time- and Size-Efficient Supercompilation, 201?
- I. Klyuchnikov and S. Romanenko. Proving the Equivalence of Higher-Order Terms by Means of Supercompilation. In A. Pnueli, I. Virbitskaite, and A. Voronkov, editors, *Perspectives of Systems Informatics*, volume 5947 of *Lecture Notes in Computer Science*, pages 193–205. Springer Berlin Heidelberg, 2010. ISBN 978-3-642-11485-4. doi: 10.1007/978-3-642-11486-1\_17. URL [http://dx.doi.org/10.1007/978-3-642-11486-1\\_17](http://dx.doi.org/10.1007/978-3-642-11486-1_17).

## References III

- J. P. Magalhães, S. Holdermans, J. Jeuring, and A. Löh. Optimizing Generics is Easy! In *Proceedings of the 2010 ACM SIGPLAN Workshop on Partial Evaluation and Program Manipulation*, PEPM '10, pages 33–42, New York, NY, USA, 2010. ACM. ISBN 978-1-60558-727-1. doi: 10.1145/1706356.1706366. URL <http://doi.acm.org/10.1145/1706356.1706366>.
- S. Peyton Jones. Call-pattern Specialisation for Haskell Programs. In *Proceedings of the 12th ACM SIGPLAN International Conference on Functional Programming*, ICFP '07, pages 327–337, New York, NY, USA, 2007. ACM. ISBN 978-1-59593-815-2. doi: 10.1145/1291151.1291200. URL <http://doi.acm.org/10.1145/1291151.1291200>.
- V. F. Turchin. The Concept of a Supercompiler. *ACM Transactions on Programming Languages and Systems*, 8:292–325, 1986.

## References IV

P. Wadler. Deforestation: Transforming Programs to Eliminate Trees. *Theor. Comput. Sci.*, 73(2):231–248, Jan. 1988. ISSN 0304-3975. doi: 10.1016/0304-3975(90)90147-A. URL [http://dx.doi.org/10.1016/0304-3975\(90\)90147-A](http://dx.doi.org/10.1016/0304-3975(90)90147-A).